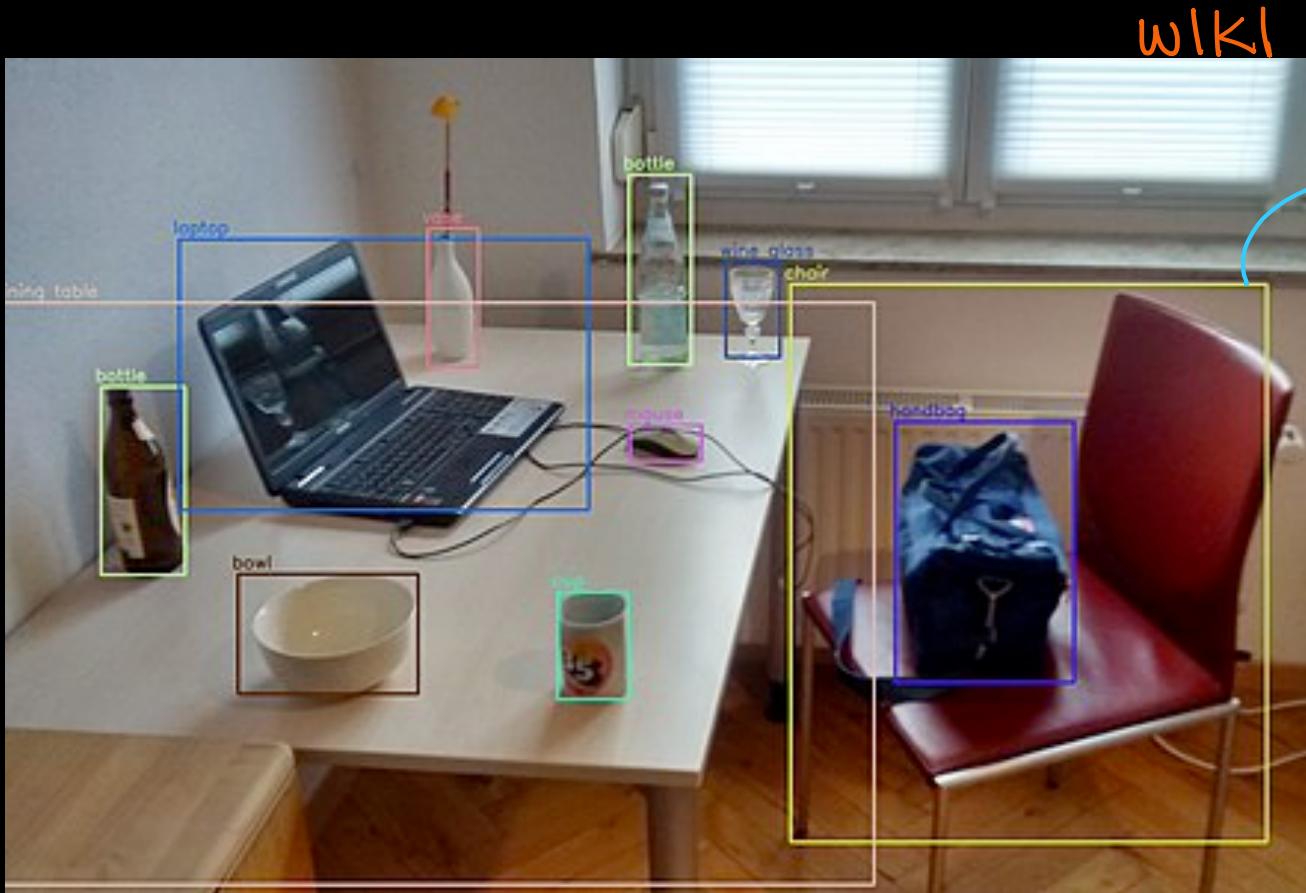


Object-detection

Agenda

1. Problem formulation
2. Azure API-based (+code)
3. Overview of Deep-Learning based models
4. YOLO v3 dive - deep
 - a. Theory
 - b. C-code

Problem-formulation:



wiki

bounding boxes
(NOT pixels)

multiple-objects in single image

Using Azure-API

- Pre-req: Using Web-APIs in Python

Overview:

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-object-detection>

Documentation:

<https://westcentralus.dev.cognitive.microsoft.com/docs/services/5adf991815e1060e6355ad44/operations/56f91f2e778daf14a499e1fa>

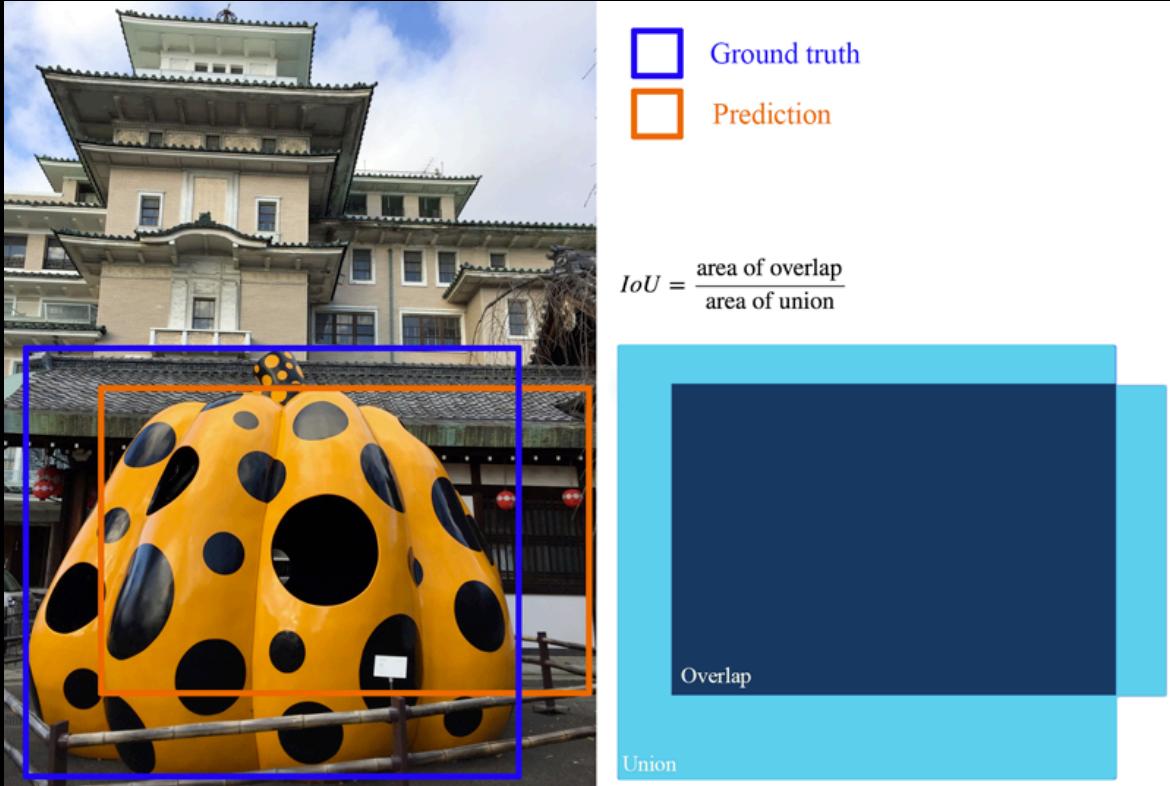
Python code:

<https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/quickstarts/python-analyze>

```
print(json.dumps(response.json()))
```

(no 'objects' in
visualFeatures)

Performance - Metric :



https://medium.com/@jonathan_hui/map-mean-average-precision-for-object-detection-45c121a31173

Prediction is
correct if

$IoU \geq 0.5$

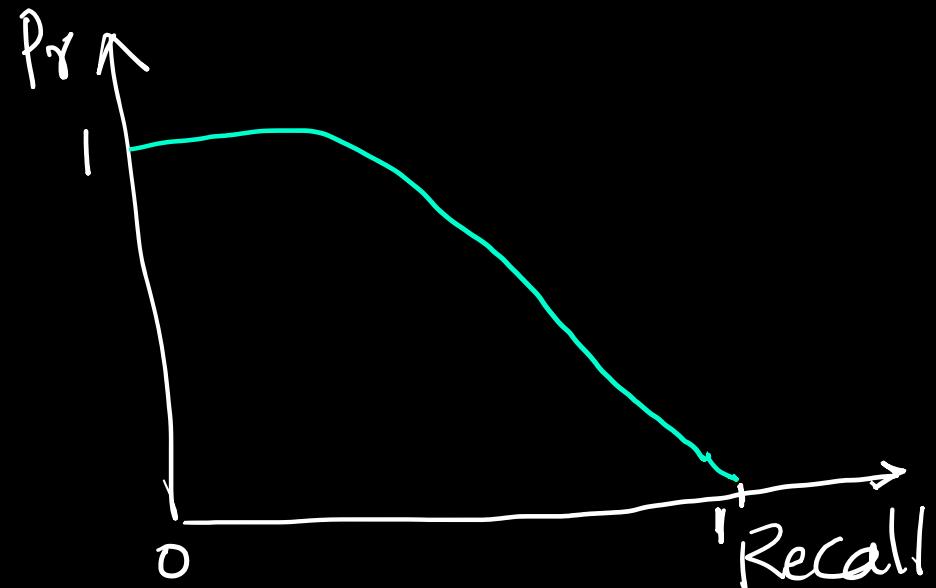
Performance - Metric :

Now, for each class (chair, person, ...),

compute avg-precision



AU-PR Curve



Performance - Metric :

mean - avg - precision (mAP) → NOT Max - A - posterior
(MAP)



mean across all classes

Deep-Learning - Models:

input : image / video

output : bounding-box-1 , object-class-1
 bounding-box-2 , object-class-2
 :
 :

Dataset : COCO <http://cocodataset.org/#home> [80-class-data]

Trade off: Speed VS mAP

→ medical diagnosis

↙
self-driving cars

real-time face-detection

R-CNN

fast R-CNN

faster R-CNN

SSD

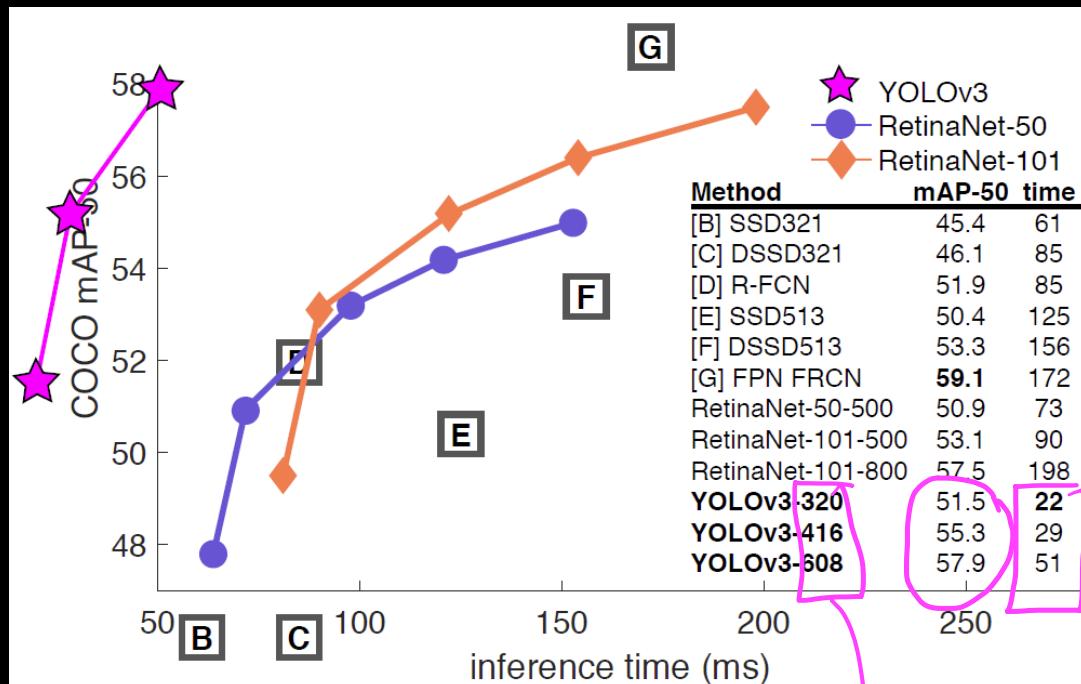
RetinaNet

YOLO V1

YOLO V2 / YOLO 9000

YOLO V3

APR-2018



input - image size ↗

Super-fast + v-good mAP

YOLO - V3: feature-extractor

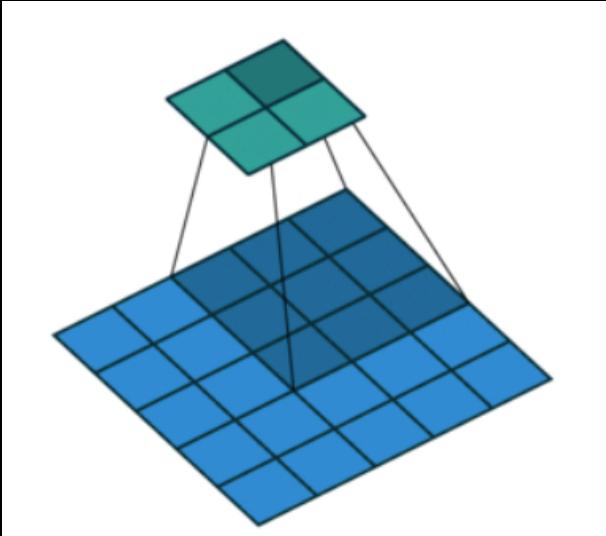
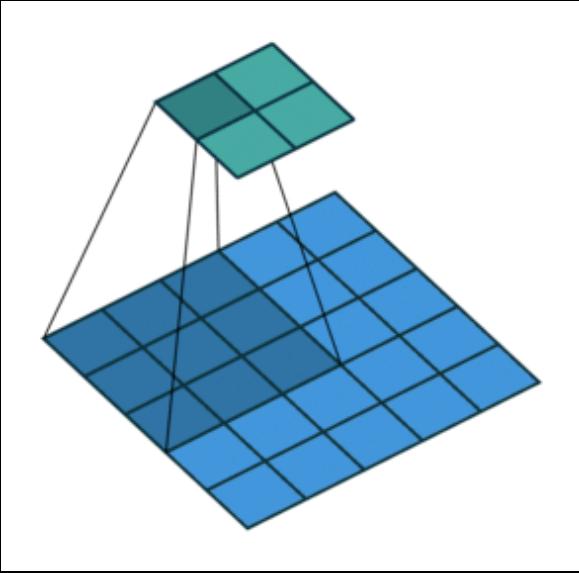
①

- fully convolutional Network (FCN)
- CONV - BatchNorm - LeakyReLU
- CONV with stride - downsample
- NO pooling
- pre-trained model

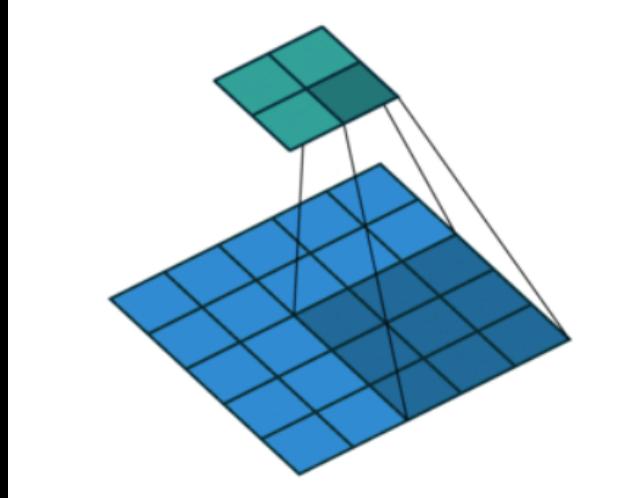
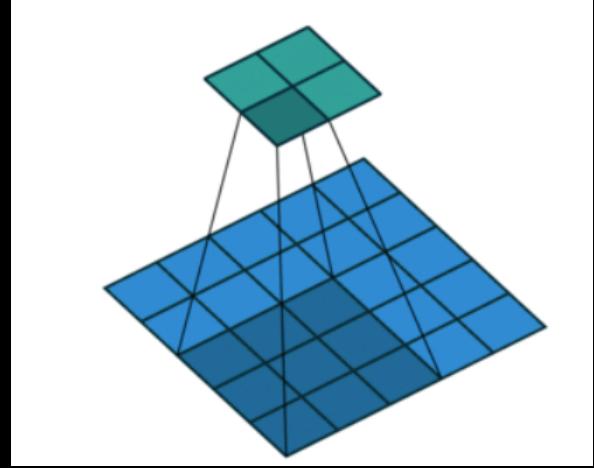
| Type | Filters | Size | Output |
|------------------|---------|------------------|------------------|
| Convolutional | 32 | 3×3 | 256×256 |
| Convolutional | 64 | $3 \times 3 / 2$ | 128×128 |
| 1x Convolutional | 32 | 1×1 | |
| 1x Convolutional | 64 | 3×3 | |
| | | Residual | 128×128 |
| Convolutional | 128 | $3 \times 3 / 2$ | 64×64 |
| 2x Convolutional | 64 | 1×1 | |
| 2x Convolutional | 128 | 3×3 | |
| | | Residual | 64×64 |
| Convolutional | 256 | $3 \times 3 / 2$ | 32×32 |
| 8x Convolutional | 128 | 1×1 | |
| 8x Convolutional | 256 | 3×3 | |
| | | Residual | 32×32 |
| Convolutional | 512 | $3 \times 3 / 2$ | 16×16 |
| 8x Convolutional | 256 | 1×1 | |
| 8x Convolutional | 512 | 3×3 | |
| | | Residual | 16×16 |
| Convolutional | 1024 | $3 \times 3 / 2$ | 8×8 |
| 4x Convolutional | 512 | 1×1 | |
| 4x Convolutional | 1024 | 3×3 | |
| | | Residual | 8×8 |
| Avgpool | | Global | |
| Connected | | 1000 | |
| Softmax | | | |

Darknet-53 model

<https://pjreddie.com/media/files/papers/YOLOv3.pdf>



Stride - 2
convolutions



5x5 $\xrightarrow{3 \times 3 \text{ kernels}} 2 \times 2$

https://github.com/vdumoulin/conv_arithmetic

a.k.a \rightarrow feature-extractor

Used
in
YOLOv2

| Backbone | Top-1 | Top-5 | Bn Ops | BFLOP/s | FPS |
|-----------------|-------------|-------------|--------|---------|-------------------------|
| Darknet-19 [15] | 74.1 | 91.8 | 7.29 | 1246 | 171 |
| ResNet-101[5] | 77.1 | 93.7 | 19.7 | 1039 | 53 |
| ResNet-152 [5] | 77.6 | 93.8 | 29.4 | 1090 | 37 |
| Darknet-53 | 77.2 | 93.8 | 18.7 | 1457 | 78 \rightarrow V.good |

Slowest (Darknet-19)
More Ops/sec

input - Images :

416 × 416 × 3

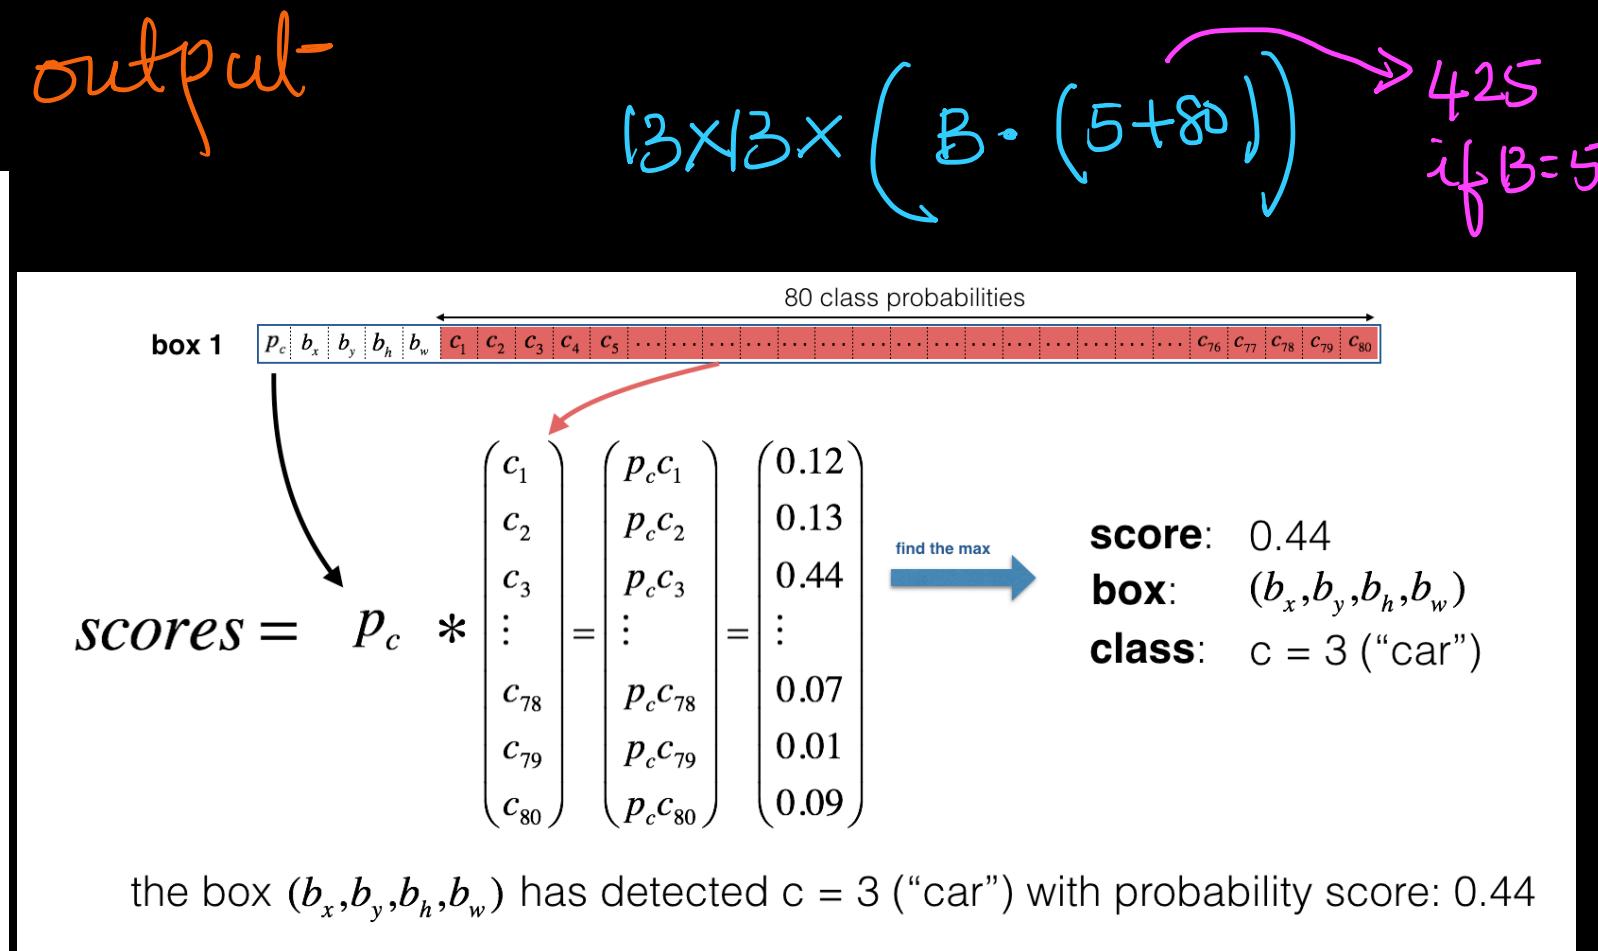
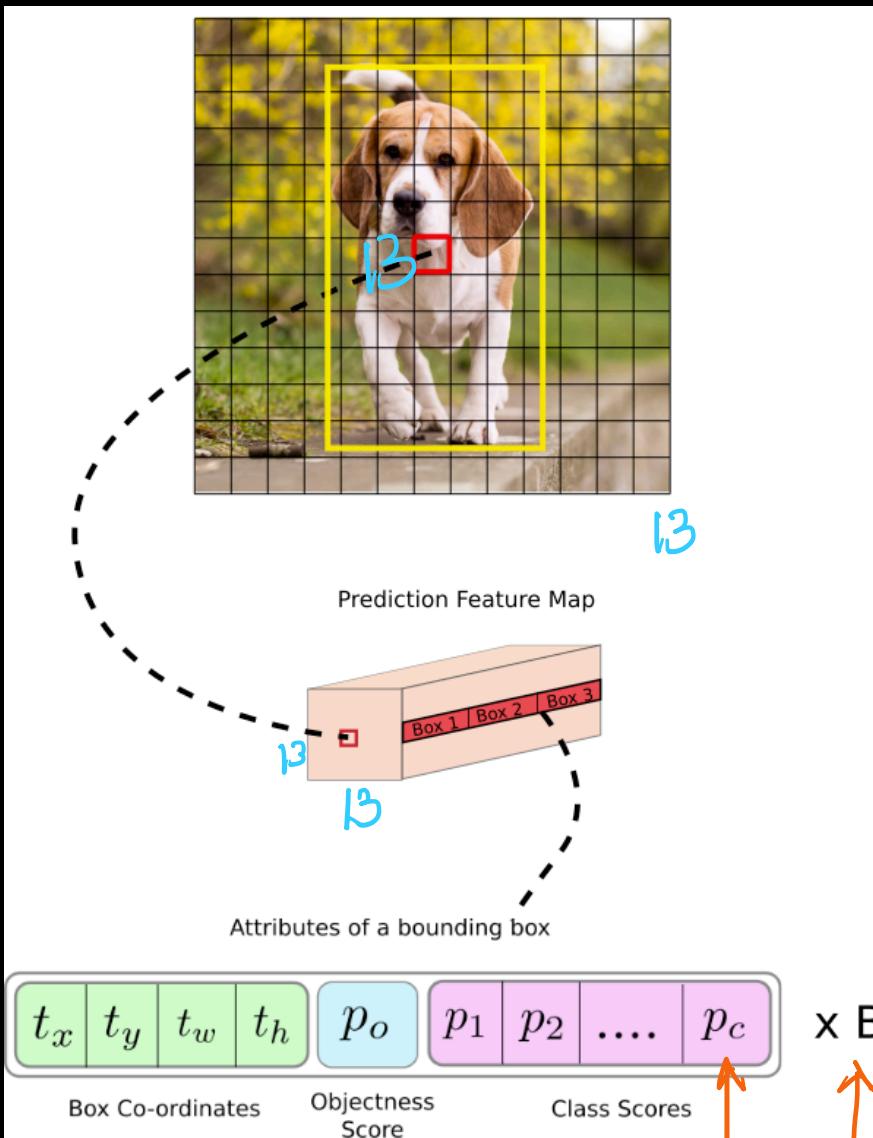
feature-extracted : 13 × 13 × 1024

| Type | Filters | Size | Output |
|------------------|---------|-----------|-----------|
| Convolutional | 32 | 3 × 3 | 256 × 256 |
| Convolutional | 64 | 3 × 3 / 2 | 128 × 128 |
| 1x Convolutional | 32 | 1 × 1 | |
| Convolutional | 64 | 3 × 3 | |
| Residual | | | 128 × 128 |
| Convolutional | 128 | 3 × 3 / 2 | 64 × 64 |
| 2x Convolutional | 64 | 1 × 1 | |
| Convolutional | 128 | 3 × 3 | |
| Residual | | | 64 × 64 |
| Convolutional | 256 | 3 × 3 / 2 | 32 × 32 |
| 8x Convolutional | 128 | 1 × 1 | |
| Convolutional | 256 | 3 × 3 | |
| Residual | | | 32 × 32 |
| Convolutional | 512 | 3 × 3 / 2 | 16 × 16 |
| 8x Convolutional | 256 | 1 × 1 | |
| Convolutional | 512 | 3 × 3 | |
| Residual | | | 16 × 16 |
| Convolutional | 1024 | 3 × 3 / 2 | 8 × 8 |
| 4x Convolutional | 512 | 1 × 1 | |
| Convolutional | 1024 | 3 × 3 | |
| Residual | | | 8 × 8 |
| Avgpool | | Global | |
| Connected | | 1000 | |
| Softmax | | | |

Darknet-53 model

/32

② Bounding-boxes & output



<https://pylessons.com/YOLOv3-introduction/>

$\times B$
 $\# \text{boxes}$
 $C=80$

feature-extractin

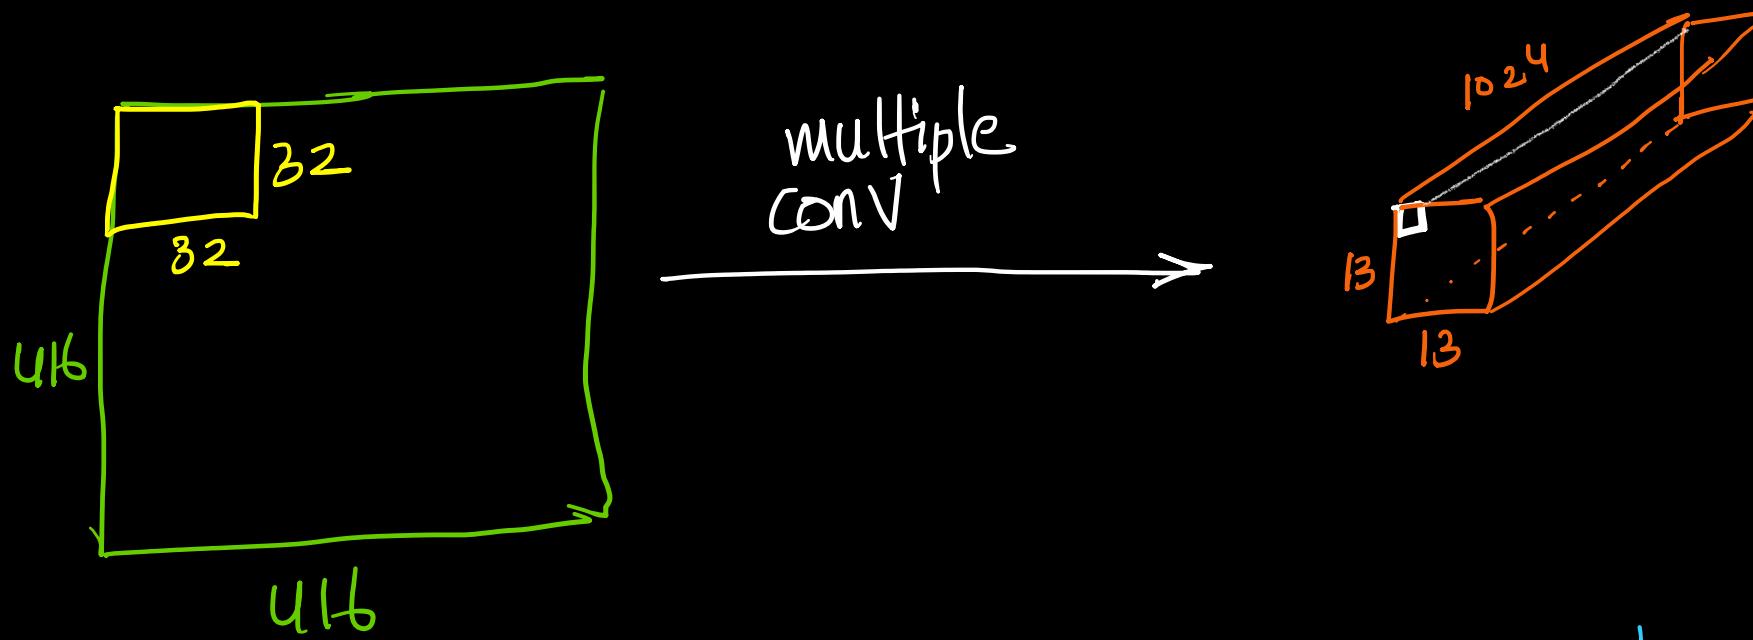
$13 \times 13 \times 1024$

$1 \times 1 \text{ conv}$
425

$13 \times 13 \times 425$

res. output

Refer: <https://www.appliedaicourse.com/lecture/11/applied-machine-learning-online-course/3412/inception-network/8/module-8-neural-networks-computer-vision-and-deep-learning>



effective receptive
field

$$\frac{416}{B} = 32$$

bounding box representation

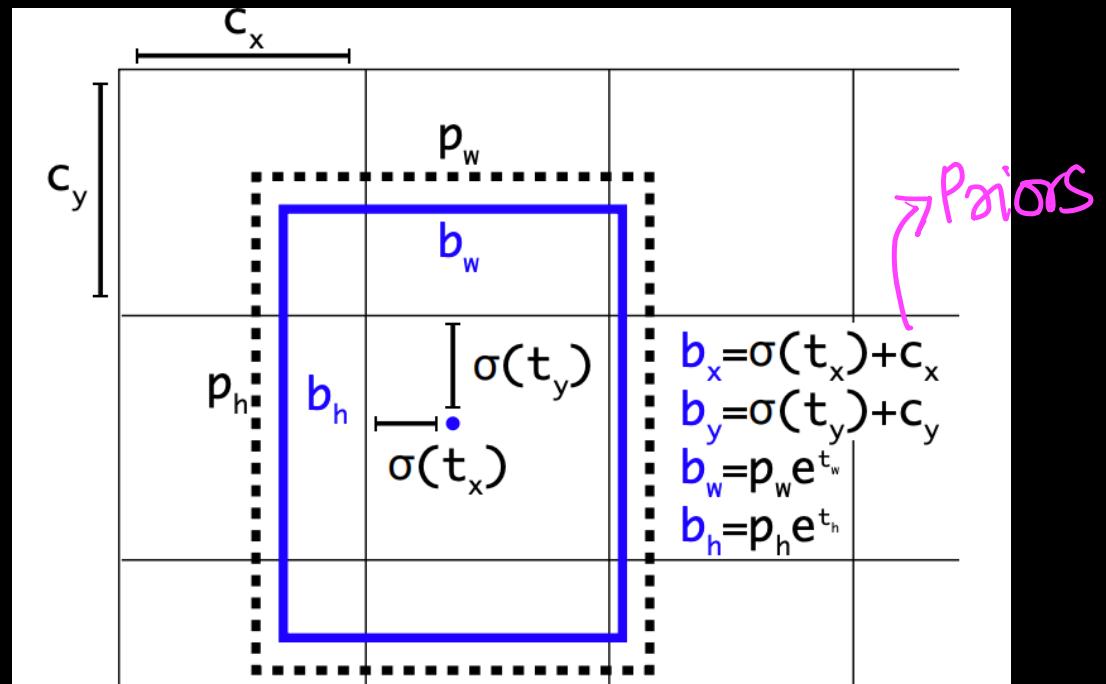
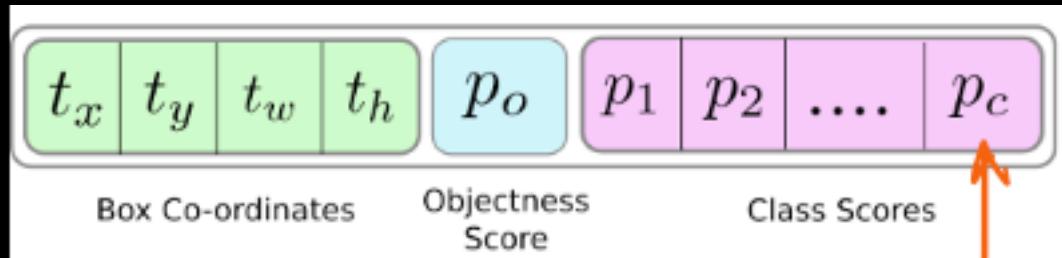


Figure 2. **Bounding boxes with dimension priors and location prediction.** We predict the width and height of the box as offsets from cluster centroids. We predict the center coordinates of the box relative to the location of filter application using a sigmoid function. This figure blatantly self-plagiarized from [15].

ANCHOR-boxes:

$$[c_x, c_y, p_w, p_h]$$

<https://pjreddie.com/media/files/papers/YOLOv3.pdf>

ANCHOR-boxes $\{x_c, y_c, P_w, P_h\}$

→ K-means clustering on train data's bounding boxes.

→ $K=5 \Rightarrow$ 5 anchor-boxes



5 bounding boxes
per cell

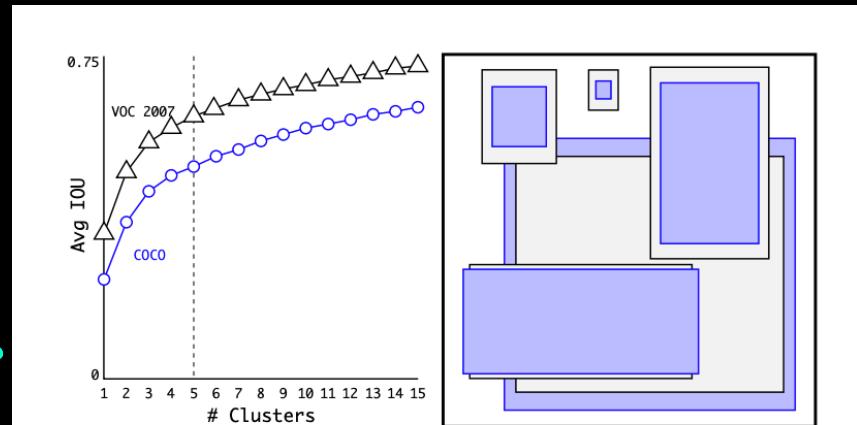
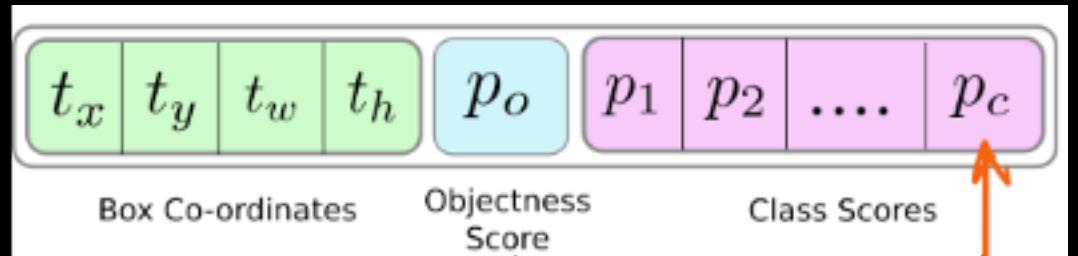


Figure 2: Clustering box dimensions on VOC and COCO. We run k-means clustering on the dimensions of bounding boxes to get good priors for our model. The left image shows the average IOU we get with various choices for k . We find that $k = 5$ gives a good tradeoff for recall vs. complexity of the model. The right image shows the relative centroids for VOC and COCO. Both sets of priors favor thinner, taller boxes while COCO has greater variation in size than VOC.

③ Per-class Sigmoids



$P(\text{box contains } C_i \text{ object})$

$$= p_o \times p_i$$

Is there any
object at all

e.g.: person, woman
 C_{10} C_{13} [Softmax
does not work]

LOSS:

$$\lambda_{\text{coord}} = 5$$

$$\lambda_{\text{noobj}} = 0.5$$

$\lambda_{\text{coord}} \cdot \text{squared-loss on}$
 f_{tw}, f_{h}, t_x, t_y

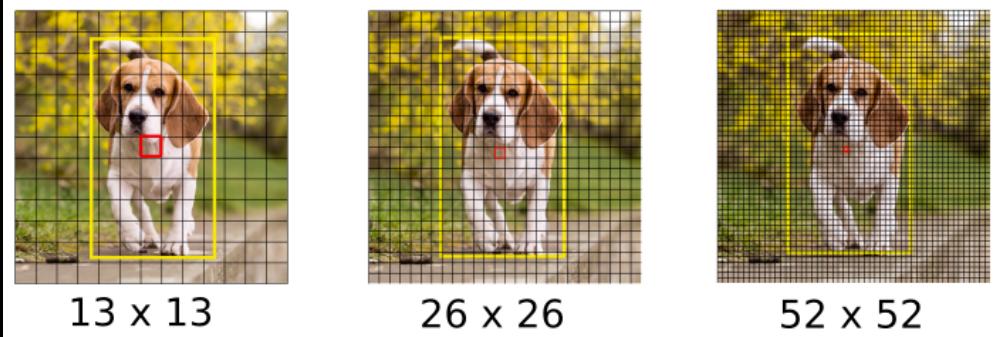
+ log-loss for p_o

all bounding
boxes
containing an object

+ log-loss for each p_i

$\sum_{\text{noobj boxes}} \lambda_{\text{noobj}} \cdot \text{log-loss}$ for p_o

④ Multi-scale Prediction

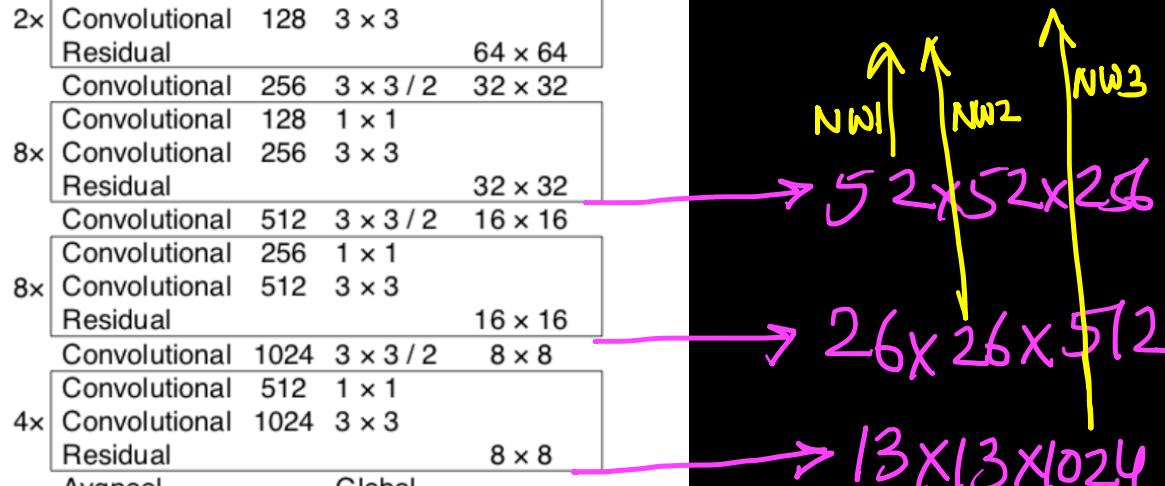


→ detect smaller objects

→ ideas from feature
pyramid networks (FPN)

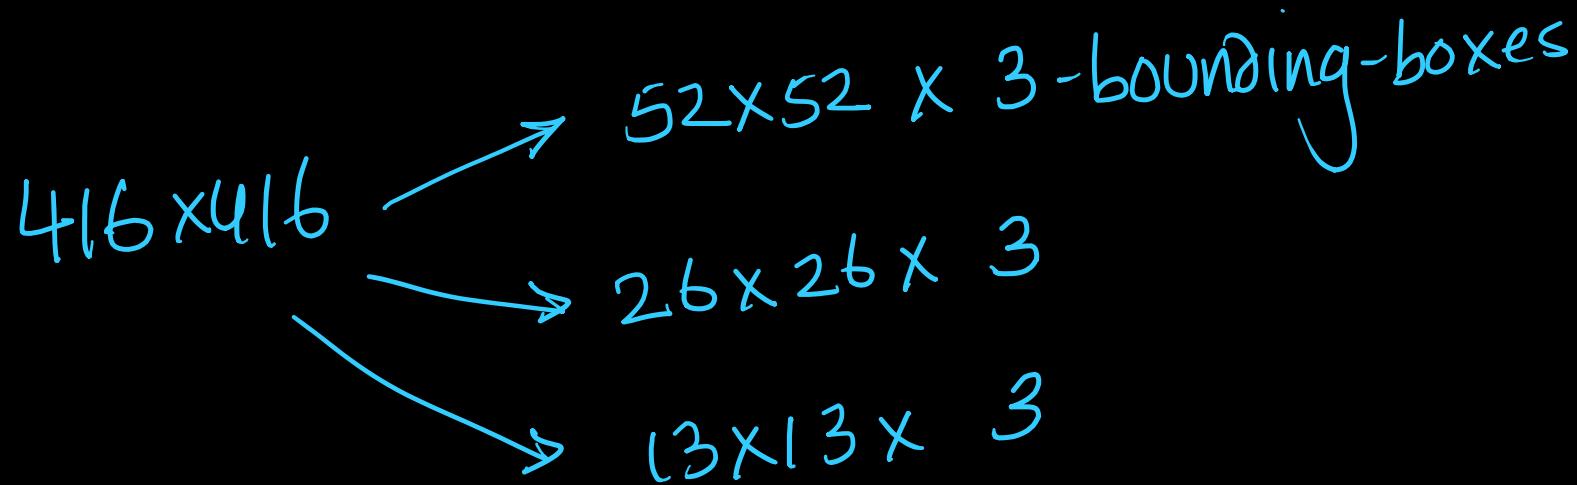
| Type | Filters | Size | Output |
|------------------|---------|------------------|-----------|
| Convolutional | 32 | 3×3 | 256 x 256 |
| Convolutional | 64 | $3 \times 3 / 2$ | 128 x 128 |
| Convolutional | 32 | 1×1 | |
| 1x Convolutional | 64 | 3×3 | 128 x 128 |
| Residual | | | |
| Convolutional | 128 | $3 \times 3 / 2$ | 64 x 64 |
| Convolutional | 64 | 1×1 | |
| 2x Convolutional | 128 | 3×3 | 64 x 64 |
| Residual | | | |
| Convolutional | 256 | $3 \times 3 / 2$ | 32 x 32 |
| Convolutional | 128 | 1×1 | |
| 8x Convolutional | 256 | 3×3 | 32 x 32 |
| Residual | | | |
| Convolutional | 512 | $3 \times 3 / 2$ | 16 x 16 |
| Convolutional | 256 | 1×1 | |
| 8x Convolutional | 512 | 3×3 | 16 x 16 |
| Residual | | | |
| Convolutional | 1024 | $3 \times 3 / 2$ | 8 x 8 |
| Convolutional | 512 | 1×1 | |
| 4x Convolutional | 1024 | 3×3 | 8 x 8 |
| Residual | | | |
| Avgpool | | Global | |
| Connected | | 1000 | |
| Softmax | | | |

Darknet-53 model



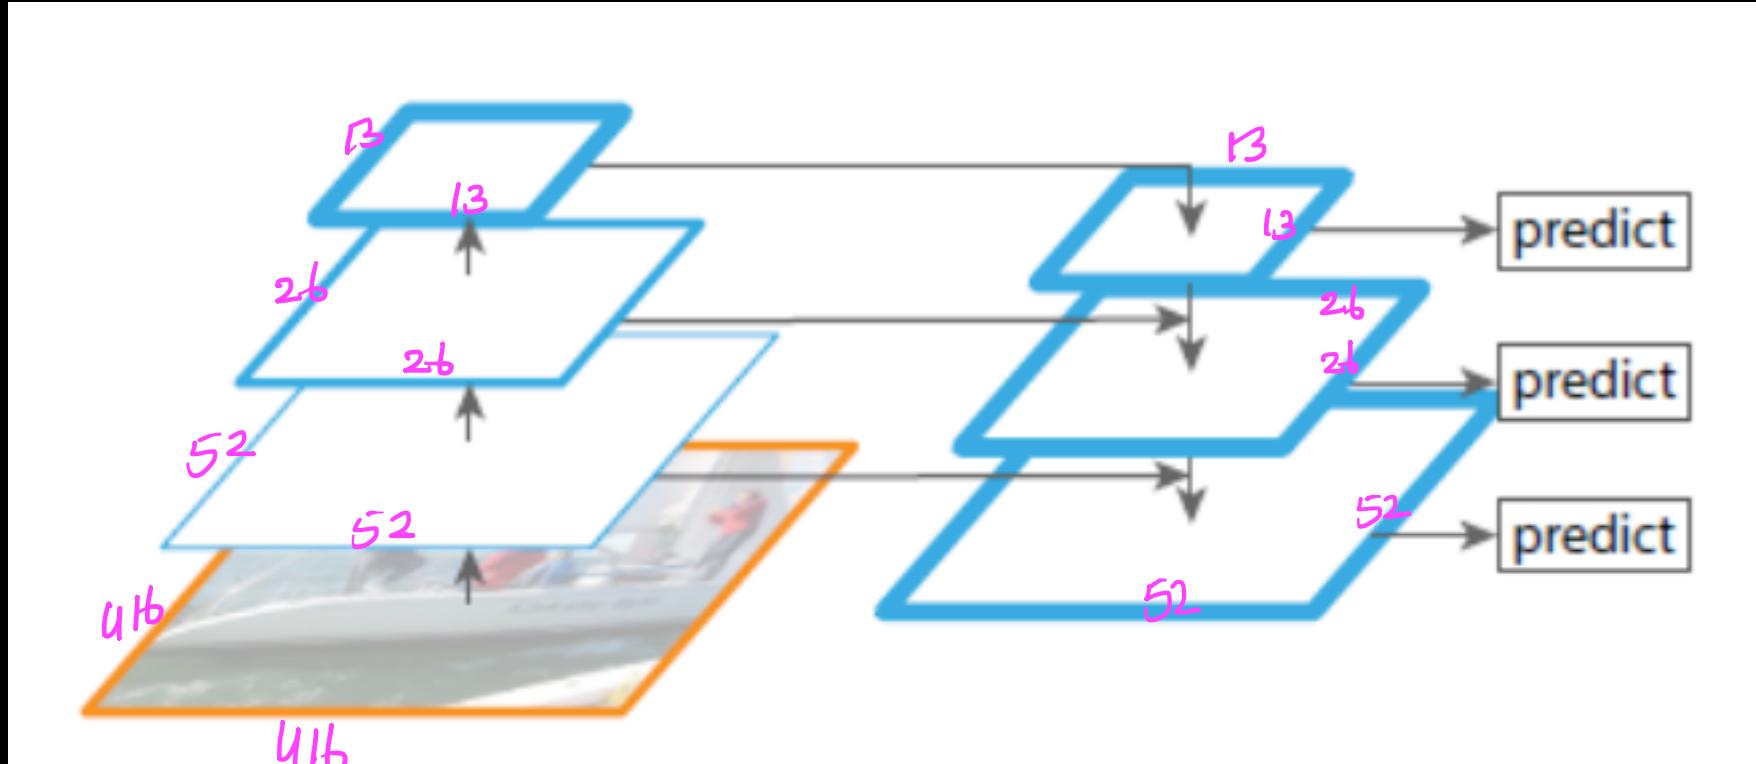
⑤

Combining boxes from various scales



10,647 boxes

FPN:



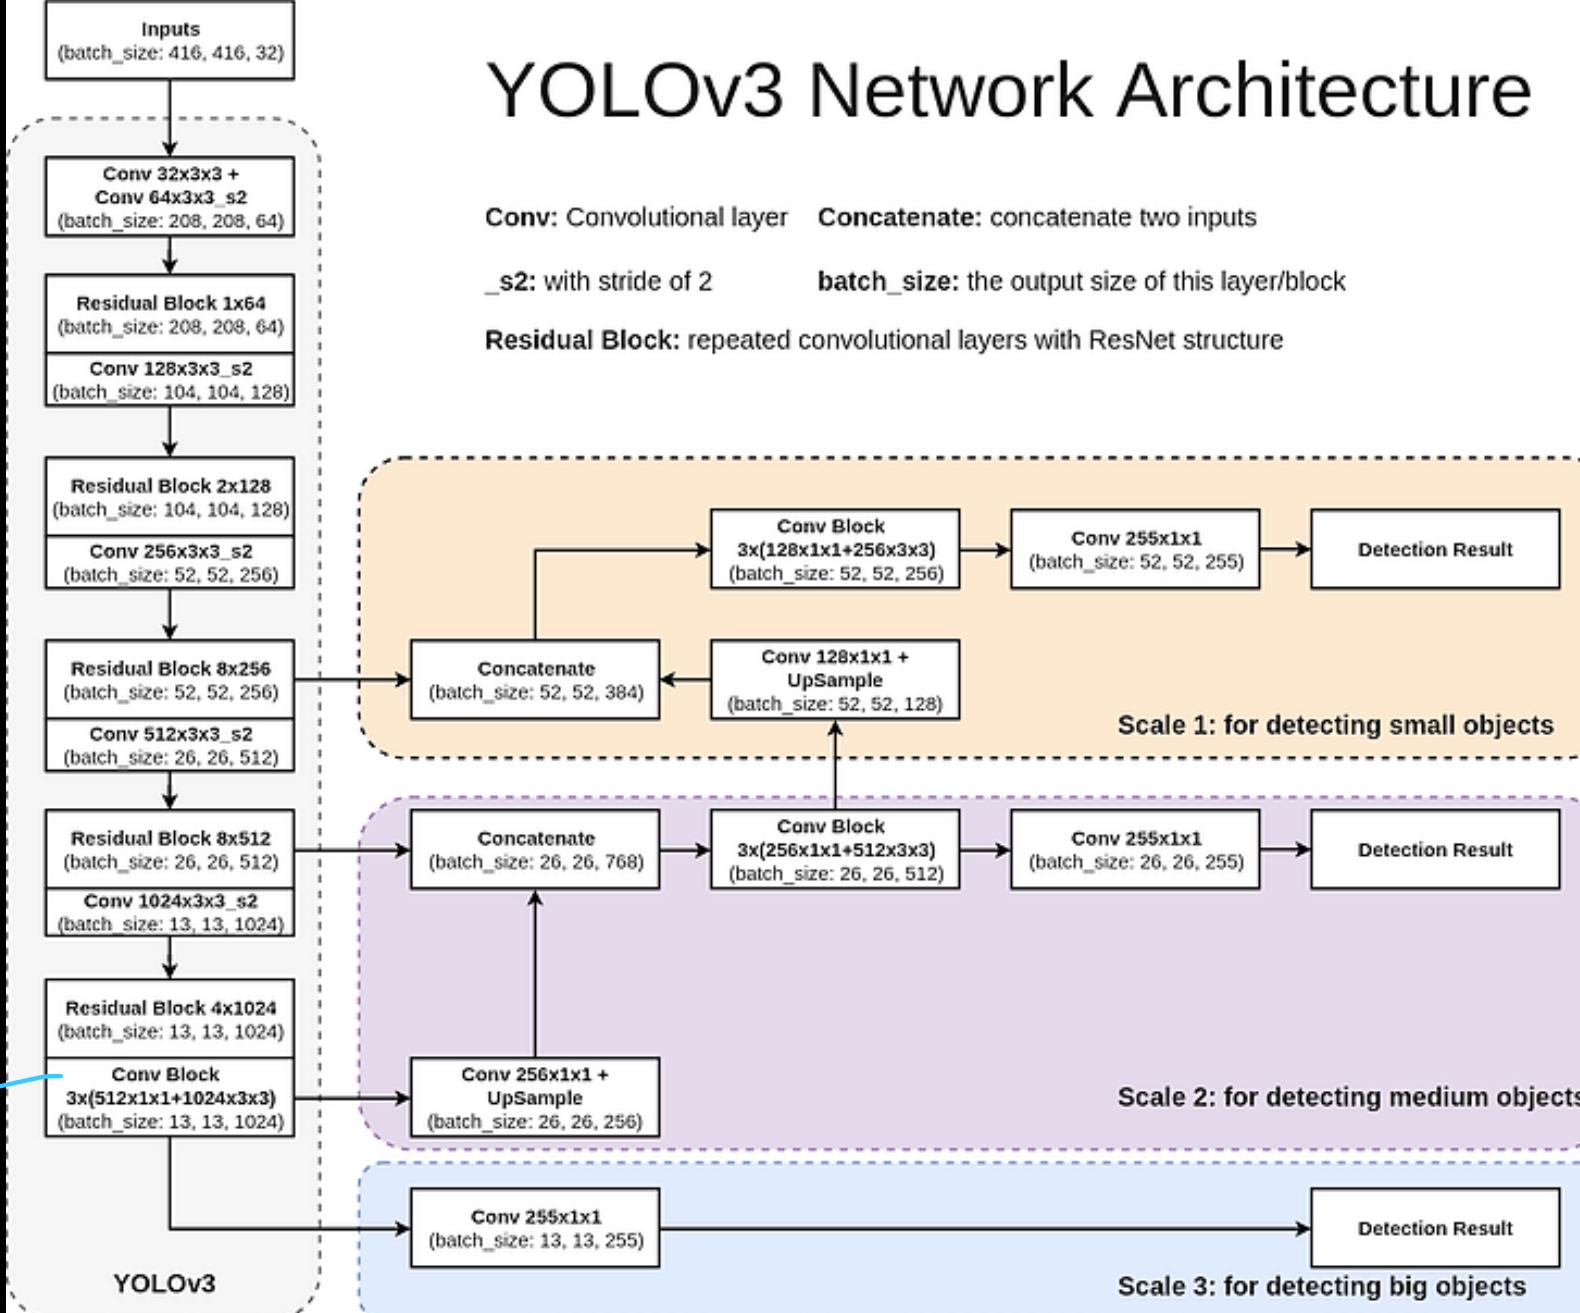
<https://towardsdatascience.com/review-fpn-feature-pyramid-network-object-detection-262fc7482610>

()

Darknet-

()

FPN with lateral-
connections



<https://www.tejashwi.io/training-yolov3-using-custom-dataset-on-google-colab/>

```

32 def make_yolov3_model():
33     input_image = Input(shape=(None, None, 3))
34     # Layer 0 => 4
35     x = _conv_block(input_image, [{"filter": 32, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 0},
36                                 {"filter": 64, "kernel": 3, "stride": 2, "bnorm": True, "leaky": True, "layer_idx": 1},
37                                 {"filter": 32, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 2},
38                                 {"filter": 64, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 3}])
39     # Layer 5 => 8
40     x = _conv_block(x, [{"filter": 128, "kernel": 3, "stride": 2, "bnorm": True, "leaky": True, "layer_idx": 5},
41                          {"filter": 64, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 6},
42                          {"filter": 128, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 7}])
43     # Layer 9 => 11
44     x = _conv_block(x, [{"filter": 64, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 9},
45                          {"filter": 128, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 10}])
46     # Layer 12 => 15
47     x = _conv_block(x, [{"filter": 256, "kernel": 3, "stride": 2, "bnorm": True, "leaky": True, "layer_idx": 12},
48                          {"filter": 128, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 13},
49                          {"filter": 256, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 14}])

```

]- repeat
2x

| Type | Filters | Size | Output |
|------|---------------|--------|------------------|
| 0 | Convolutional | 32 | 256×256 |
| 1 | Convolutional | 64 | 128×128 |
| 2 | Convolutional | 32 | 1×1 |
| 3 | Convolutional | 64 | 3×3 |
| | Residual | | 128×128 |
| 5 | Convolutional | 128 | $3 \times 3 / 2$ |
| 6 | Convolutional | 64 | 1×1 |
| 7 | 2x | 128 | 3×3 |
| | Residual | | 64×64 |
| 12 | Convolutional | 256 | $3 \times 3 / 2$ |
| B | Convolutional | 128 | 1×1 |
| 8x | Convolutional | 256 | 3×3 |
| H | Residual | | 32×32 |
| | Convolutional | 512 | $3 \times 3 / 2$ |
| | Convolutional | 256 | 1×1 |
| 8x | Convolutional | 512 | 3×3 |
| | Residual | | 16×16 |
| | Convolutional | 1024 | $3 \times 3 / 2$ |
| | Convolutional | 512 | 1×1 |
| 4x | Convolutional | 1024 | 3×3 |
| | Residual | | 8×8 |
| | Avgpool | Global | |
| | Connected | 1000 | |
| | Softmax | | |

512x512x3
26x26x512
13x13x1024

<https://machinelearningmastery.com/how-to-perform-object-detection-with-yolov3-in-keras/>

| Type | Filters | Size | Output |
|---------------|---------------|------------------|------------------|
| Convolutional | 32 | 3×3 | 256×256 |
| Convolutional | 64 | $3 \times 3 / 2$ | 128×128 |
| 1x | Convolutional | 32 | 1×1 |
| Convolutional | 64 | 3×3 | 128×128 |
| Residual | | | 128×128 |
| Convolutional | 128 | $3 \times 3 / 2$ | 64×64 |
| Convolutional | 64 | 1×1 | |
| 2x | Convolutional | 128 | 3×3 |
| Convolutional | Residual | | 64×64 |
| Convolutional | 256 | $3 \times 3 / 2$ | 32×32 |
| Convolutional | 128 | 1×1 | |
| 8x | Convolutional | 256 | 3×3 |
| Convolutional | Residual | | 32×32 |
| Convolutional | 512 | $3 \times 3 / 2$ | 16×16 |
| Convolutional | 256 | 1×1 | |
| 8x | Convolutional | 512 | 3×3 |
| Convolutional | Residual | | 16×16 |
| Convolutional | 1024 | $3 \times 3 / 2$ | 8×8 |
| Convolutional | 512 | 1×1 | |
| 4x | Convolutional | 1024 | 3×3 |
| Convolutional | Residual | | 8×8 |
| Avgpool | | Global | |
| Connected | | 1000 | |
| Softmax | | | |

Darknet-53 model

$512 \times 512 \times 256$

$26 \times 26 \times 512$

$13 \times 13 \times 1024$

+ more times

```
# Layer 16 => 36
for i in range(7):
    x = _conv_block(x, [{"filter": 128, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 16+i*3},
                         {"filter": 256, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 17+i*3}])
skip_36 = x
# Layer 37 => 40
x = _conv_block(x, [{"filter": 512, "kernel": 3, "stride": 2, "bnorm": True, "leaky": True, "layer_idx": 37},
                     {"filter": 256, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 38},
                     {"filter": 512, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 39}])
```

<https://machinelearningmastery.com/how-to-perform-object-detection-with-yolov3-in-keras/>

```

# Layer 66 => 74
for i in range(3):
    x = _conv_block(x, [{"filter": 512, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 66+i*3},
                        {"filter": 1024, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 67+i*3}])
# Layer 75 => 79
x = _conv_block(x, [{"filter": 512, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 75},
                     {"filter": 1024, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 76},
                     {"filter": 512, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 77},
                     {"filter": 1024, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 78},
                     {"filter": 512, "kernel": 1, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 79}], skip=False)
# Layer 80 => 82
yolo_82 = _conv_block(x, [{"filter": 1024, "kernel": 3, "stride": 1, "bnorm": True, "leaky": True, "layer_idx": 80},
                           {"filter": 255, "kernel": 1, "stride": 1, "bnorm": False, "leaky": False, "layer_idx": 81}], skip=False)

```

| Type | Filters | Size | Output |
|---------------|---------------|------------------|------------------|
| Convolutional | 32 | 3×3 | 256×256 |
| Convolutional | 64 | $3 \times 3 / 2$ | 128×128 |
| 1x | 32 | 1×1 | |
| Convolutional | 64 | 3×3 | |
| | Residual | | 128×128 |
| | Convolutional | $128 / 2$ | 64×64 |
| 2x | 64 | 1×1 | |
| Convolutional | 128 | 3×3 | |
| | Residual | | 64×64 |
| | Convolutional | $256 / 2$ | 32×32 |
| 8x | 128 | 1×1 | |
| Convolutional | 256 | 3×3 | |
| | Residual | | 32×32 |
| | Convolutional | $512 / 2$ | 16×16 |
| 8x | 256 | 1×1 | |
| Convolutional | 512 | 3×3 | |
| | Residual | | 16×16 |
| | Convolutional | $1024 / 2$ | 8×8 |
| 4x | 512 | 1×1 | |
| Convolutional | 1024 | 3×3 | |
| | Residual | | 8×8 |
| | Avgpool | Global | |
| | Connected | 1000 | |
| | Softmax | | |

Darknet-53 model

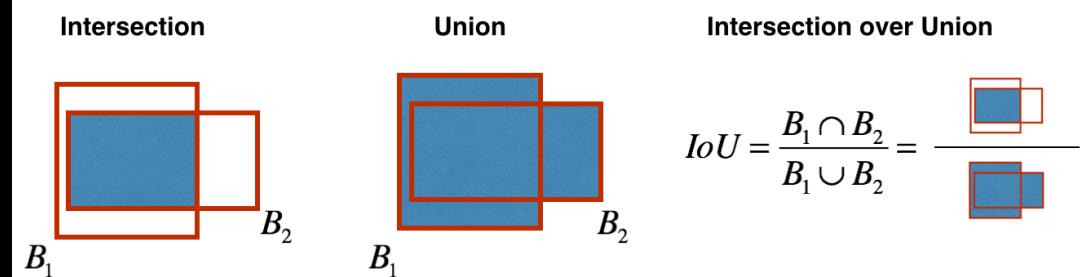
<https://machinelearningmastery.com/how-to-perform-object-detection-with-yolov3-in-keras/>



a) Filter: boxes with $\max_i(p_o \cdot p_i) < 0.5$ are ignored

b) Non-Max-Suppression (NMS)

Pick the box
with Max IoU



Code :

<https://pjreddie.com/darknet/yolo/> → fast C-implementation

Keras-implementation : <https://github.com/qwwweee/keras-yolo3>
↳ not as-fast

<https://machinelearningmastery.com/how-to-perform-object-detection-with-yolov3-in-keras/>